





Same 14 tasks and model; one fresh agent per run; every run passed 14/14 (verified by the benchmark server).

arc-browser bars: mean of its two final runs (v10, v11), whiskers show their range. ego-lite and Claude for Chrome: one run each, stock setup. arc-browser was tuned on these tasks, so the gap is an upper bound.