




Same 14 tasks and model; one fresh agent per run; every run passed 14/14 (verified by the benchmark server).

claude4arc: mean of two runs, whiskers show their range. ego-lite and Claude for Chrome: one run each, stock setup.

claude4arc was tuned on these tasks, so the gap is an upper bound.